

# graphsage-final

March 26, 2024

```
[1]: import torch
import numpy as np
import pandas as pd
import pyarrow.parquet as pq
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader, random_split
import torch.nn.functional as F
from torchvision import models

import torch.optim as optim
from tqdm import tqdm

from sklearn.metrics import roc_auc_score, confusion_matrix, roc_curve
import seaborn as sns
from skimage.segmentation import slic
```

```
[2]: !pip install torch_geometric
!pip install networkx
```

```
/opt/conda/lib/python3.10/pty.py:89: RuntimeWarning: os.fork() was called.
os.fork() is incompatible with multithreaded code, and JAX is multithreaded, so
this will likely lead to a deadlock.
```

```
pid, fd = os.forkpty()
```

```
Collecting torch_geometric
```

```
Downloading torch_geometric-2.5.2-py3-none-any.whl.metadata (64 kB)
```

```
64.2/64.2 kB
```

```
3.1 MB/s eta 0:00:00
```

```
Requirement already satisfied: tqdm in /opt/conda/lib/python3.10/site-
packages (from torch_geometric) (4.66.1)
```

```
Requirement already satisfied: numpy in /opt/conda/lib/python3.10/site-packages
(from torch_geometric) (1.26.4)
```

```
Requirement already satisfied: scipy in /opt/conda/lib/python3.10/site-packages
(from torch_geometric) (1.11.4)
```

Requirement already satisfied: fsspec in /opt/conda/lib/python3.10/site-packages (from torch\_geometric) (2024.3.0)

Requirement already satisfied: Jinja2 in /opt/conda/lib/python3.10/site-packages (from torch\_geometric) (3.1.2)

Requirement already satisfied: aiohttp in /opt/conda/lib/python3.10/site-packages (from torch\_geometric) (3.9.1)

Requirement already satisfied: requests in /opt/conda/lib/python3.10/site-packages (from torch\_geometric) (2.31.0)

Requirement already satisfied: pyparsing in /opt/conda/lib/python3.10/site-packages (from torch\_geometric) (3.1.1)

Requirement already satisfied: scikit-learn in /opt/conda/lib/python3.10/site-packages (from torch\_geometric) (1.2.2)

Requirement already satisfied: psutil>=5.8.0 in /opt/conda/lib/python3.10/site-packages (from torch\_geometric) (5.9.3)

Requirement already satisfied: attrs>=17.3.0 in /opt/conda/lib/python3.10/site-packages (from aiohttp->torch\_geometric) (23.2.0)

Requirement already satisfied: multidict<7.0,>=4.5 in /opt/conda/lib/python3.10/site-packages (from aiohttp->torch\_geometric) (6.0.4)

Requirement already satisfied: yarl<2.0,>=1.0 in /opt/conda/lib/python3.10/site-packages (from aiohttp->torch\_geometric) (1.9.3)

Requirement already satisfied: frozenlist>=1.1.1 in /opt/conda/lib/python3.10/site-packages (from aiohttp->torch\_geometric) (1.4.1)

Requirement already satisfied: aiosignal>=1.1.2 in /opt/conda/lib/python3.10/site-packages (from aiohttp->torch\_geometric) (1.3.1)

Requirement already satisfied: async-timeout<5.0,>=4.0 in /opt/conda/lib/python3.10/site-packages (from aiohttp->torch\_geometric) (4.0.3)

Requirement already satisfied: MarkupSafe>=2.0 in /opt/conda/lib/python3.10/site-packages (from Jinja2->torch\_geometric) (2.1.3)

Requirement already satisfied: charset-normalizer<4,>=2 in /opt/conda/lib/python3.10/site-packages (from requests->torch\_geometric) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in /opt/conda/lib/python3.10/site-packages (from requests->torch\_geometric) (3.6)

Requirement already satisfied: urllib3<3,>=1.21.1 in /opt/conda/lib/python3.10/site-packages (from requests->torch\_geometric) (1.26.18)

Requirement already satisfied: certifi>=2017.4.17 in /opt/conda/lib/python3.10/site-packages (from requests->torch\_geometric) (2024.2.2)

Requirement already satisfied: joblib>=1.1.1 in /opt/conda/lib/python3.10/site-packages (from scikit-learn->torch\_geometric) (1.3.2)

Requirement already satisfied: threadpoolctl>=2.0.0 in /opt/conda/lib/python3.10/site-packages (from scikit-learn->torch\_geometric) (3.2.0)

Downloading torch\_geometric-2.5.2-py3-none-any.whl (1.1 MB)

1.1/1.1 MB

28.6 MB/s eta 0:00:00

Installing collected packages: torch\_geometric

Successfully installed torch\_geometric-2.5.2

Requirement already satisfied: networkx in /opt/conda/lib/python3.10/site-packages (3.2.1)

```
[3]: from torch_geometric.data import Data
from sklearn.neighbors import kneighbors_graph
from torch_geometric.data import Dataset, Data, DataLoader
from sklearn.metrics import mean_squared_error
from sklearn.neighbors import NearestNeighbors
import networkx as nx
from torch_geometric.utils import to_networkx

from torch_geometric.loader import DataLoader
from torch_geometric.nn import GCNConv, global_mean_pool
from torch.nn import Linear
import torch.nn.functional as F
from torch.utils.data import random_split
from torch.nn import Sequential, Linear, ReLU
from torch_geometric.nn import NNConv, global_mean_pool
from torch_geometric.nn import GATConv
from skimage import io
```

```
[4]: chunk_size = 25

# List of Parquet file paths
file_paths = [
    '/kaggle/input/task2-24/QCDToGGQQ_IMGjet_RH1all_jet0_run0_n36272.test.
↳snappy.parquet',
    '/kaggle/input/task2-24/QCDToGGQQ_IMGjet_RH1all_jet0_run1_n47540.test.
↳snappy.parquet',
    '/kaggle/input/task2-24/QCDToGGQQ_IMGjet_RH1all_jet0_run2_n55494.test.
↳snappy.parquet'
]

# Initialize an empty list to store dataframes
dfs = []

# Loop through each file path
for file_path in file_paths:
    # Create a Parquet file reader object
    parquet_file = pq.ParquetFile(file_path)

    # Determine the total number of rows in the file
    total_rows = parquet_file.metadata.num_rows

    # Calculate the number of chunks
```

```

num_chunks = total_rows // chunk_size + (1 if total_rows % chunk_size else 0)

# Loop over the file in chunks
for chunk_index in range(num_chunks):
    # Read a chunk of rows from the file
    chunk = parquet_file.read_row_group(chunk_index, columns=None)
    df = chunk.to_pandas()

    # Append the DataFrame to the list
    dfs.append(df)

# Concatenate all the DataFrames into a single DataFrame
data = pd.concat(dfs, ignore_index=True)

```

```

[5]: def to_3d(arr):
    vishak=[]
    for i in range (0,3):
        vis=np.stack(np.stack(arr)[i],axis=-1)
        vishak.append(vis)
    vishak=np.array(vishak)
    vishak_max = vishak.max()
    vishak_min = vishak.min()
    vishak = (vishak - vishak_min)/(vishak_max - vishak_min)
    return vishak

```

```

[6]: data["X_jets"] = data["X_jets"].apply(to_3d)

```

```

[7]: def image_to_graph(image, patch_size=25, n_neighbors=15):
    """
    Convert an image to a graph of its 5x5 patches.

    Parameters:
    - image: A (125, 125, 3) numpy array.
    - patch_size: Size of the square patches (default 5).
    - n_neighbors: Number of neighbors for KNN (default 5).

    Returns:
    - nodes: An array of node features.
    - edges: A list of tuples (i, j, mse) representing edges and their MSE.
    """
    # Validate image shape

    assert image.shape[0] == image.shape[1], "Image must be square."

    # Number of patches along one dimension
    num_patches = image.shape[0] // patch_size

```

```

# Initialize nodes and edges
nodes = []
edges = []

# Create patches and flatten them to create node features
for i in range(0, image.shape[0], patch_size):
    for j in range(0, image.shape[1], patch_size):
        patch = image[i:i+patch_size, j:j+patch_size, :].reshape(-1)
        nodes.append(patch)

nodes = np.array(nodes)

# Use KNN to find nearest neighbors for each node
nbrs = NearestNeighbors(n_neighbors=n_neighbors+1,
    ↪algorithm='ball_tree').fit(nodes)
distances, indices = nbrs.kneighbors(nodes)

# Calculate MSE for each pair of neighbors and create edges
for i in range(indices.shape[0]):
    for j in range(1, indices.shape[1]): # Start from 1 to skip
    ↪self-connection
        mse = mean_squared_error(nodes[i], nodes[indices[i, j]])
        edges.append((i, indices[i, j], mse))

return nodes, edges

```

```

[8]: class QuarkGluonDataset(Dataset):

    def __init__(self, dataframe, root='', transform=None, pre_transform=None):
        """
        Custom dataset for quarks and gluons classification.

        Parameters:
        - image_list: A list of (125, 125, 3) numpy arrays.
        - labels: A list of integers (0 or 1) representing the class labels for
    ↪the images.
        """
        self.dataframe = dataframe
        super(QuarkGluonDataset, self).__init__(root, transform, pre_transform)

    def len(self):
        return len(self.dataframe)

    def get(self, idx):
        # Convert an image to graph data

```

```

        image = self.dataframe.iloc[idx]['X_jets']
        image = image.transpose(1,2,0)
        label = self.dataframe.iloc[idx]['y']
#         print(type(image))
        nodes, edges = image_to_graph(image)

        # Convert to PyTorch tensors
        x = torch.tensor(nodes, dtype=torch.float) # Node features
        edge_index = torch.tensor([(i, j) for i, j, _ in edges], dtype=torch.
→long).t().contiguous() # Edge indices
        edge_attr = torch.tensor([mse for _, _, mse in edges], dtype=torch.
→float).unsqueeze(1) # Edge attributes
        y = torch.tensor([label], dtype=torch.long) # Label

        return Data(x=x, edge_index=edge_index, edge_attr=edge_attr, y=y)

```

```

[9]: dataset = QuarkGluonDataset(data)
dataset_size = len(dataset)
train_size = int(0.8 * dataset_size)
test_size = dataset_size - train_size

# Perform the random split
train_dataset, test_dataset = random_split(dataset, [train_size, test_size])

# Create the DataLoaders for the train and test sets
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
valid_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)

```

```

[10]: from torch_geometric.nn import SAGEConv, global_mean_pool
from torch.nn import Linear, Dropout
import torch.nn.functional as F

class GraphSAGENet(torch.nn.Module):
    def __init__(self, num_node_features, num_classes):
        super(GraphSAGENet, self).__init__()
        self.sage1 = SAGEConv(num_node_features, 512)
        self.sage2 = SAGEConv(512, 256)
        self.sage3 = SAGEConv(256, 128)
        self.dropout = Dropout(0.5)
        self.lin = Linear(128, num_classes)

    def forward(self, data):
        x, edge_index, batch = data.x, data.edge_index, data.batch

        x = F.relu(self.sage1(x, edge_index))
        x = self.dropout(x)
        x = F.relu(self.sage2(x, edge_index))

```

```

        x = self.dropout(x)
        x = F.relu(self.sage3(x, edge_index))

        x = global_mean_pool(x, batch) # Pooling to use graph-level features
        x = self.dropout(x)
        x = self.lin(x)

        return F.log_softmax(x, dim=1)

# Initialize the GraphSAGE model
num_node_features = 1875 # Adjust according to your dataset
num_classes = 2 # Assuming binary classification

model = GraphSAGENet(num_node_features=num_node_features,
    ↪ num_classes=num_classes)
print(model)

```

```

GraphSAGENet(
  (sage1): SAGEConv(1875, 512, aggr=mean)
  (sage2): SAGEConv(512, 256, aggr=mean)
  (sage3): SAGEConv(256, 128, aggr=mean)
  (dropout): Dropout(p=0.5, inplace=False)
  (lin): Linear(in_features=128, out_features=2, bias=True)
)

```

```

[11]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f'Using device: {device}')
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
criterion = torch.nn.CrossEntropyLoss()
# Move the model to the chosen device
model.to(device)

# Modify the train function to include loss in tqdm
def train(model, train_loader, optimizer, criterion):
    model.train()
    total_loss = 0
    with tqdm(train_loader, desc="Training") as tepoch:
        for data in tepoch:
            data = data.to(device)
            optimizer.zero_grad()
            # Ensure the model forward call matches the expected arguments
            out = model(data)
            loss = criterion(out, data.y)
            loss.backward()
            optimizer.step()
            total_loss += loss.item() * data.num_graphs
            tepoch.set_postfix(loss=loss.item())

```

```

    return total_loss / len(train_loader.dataset)

# Similar modifications for the test/validation function, including data
→transfer to the device
def test(model, loader, criterion):
    model.eval()
    correct = 0
    total_loss = 0
    with torch.no_grad(), tqdm(loader, desc="Evaluating", leave=False) as te:
        te.epoch:
            for data in te:
                data = data.to(device)
                # Match the model forward call with expected arguments
                out = model(data)
                loss = criterion(out, data.y)
                total_loss += loss.item() * data.num_graphs
                pred = out.argmax(dim=1)
                correct += int((pred == data.y).sum())
    return correct / len(loader.dataset), total_loss / len(loader.dataset)

best_val_acc = 0.0
for epoch in range(1, 100):
    train_loss = train(model, train_loader, optimizer, criterion)
    # train_acc, _ = test(model, train_loader, criterion)
    val_acc, val_loss = test(model, valid_loader, criterion)
    if (val_loss - train_loss > 0.3):
        print("Early Stopping")
        break
    if val_acc > best_val_acc:
        best_val_acc = val_acc
        torch.save(model.state_dict(), 'best_model.pth')
        print(f"Saved Best Model: Epoch {epoch}, Val. Acc.: {val_acc:.4f}")

    print(f'Epoch: {epoch:03d}, Train Loss: {train_loss:.4f}, Val. Loss: te
    te.val_loss:.4f, Val. Acc.: {val_acc:.4f}')

```

Using device: cpu

Training: 100% | 140/140 [06:43<00:00, 2.88s/it, loss=0.588]

Saved Best Model: Epoch 1, Val. Acc.: 0.6063

Epoch: 001, Train Loss: 0.6677, Val. Loss: 0.6693, Val. Acc.: 0.6063

Training: 100% | 140/140 [06:43<00:00, 2.88s/it, loss=0.491]

Saved Best Model: Epoch 2, Val. Acc.: 0.6117

Epoch: 002, Train Loss: 0.6564, Val. Loss: 0.6668, Val. Acc.: 0.6117

Training: 100% | 140/140 [06:43<00:00, 2.88s/it, loss=0.591]



Epoch: 003, Train Loss: 0.6535, Val. Loss: 0.6678, Val. Acc.: 0.6009  
 Training: 100%| | 140/140 [06:45<00:00, 2.90s/it, loss=0.456]  
 Saved Best Model: Epoch 4, Val. Acc.: 0.6152  
 Epoch: 004, Train Loss: 0.6518, Val. Loss: 0.6693, Val. Acc.: 0.6152  
 Training: 100%| | 140/140 [06:46<00:00, 2.90s/it, loss=0.61]  
 Epoch: 005, Train Loss: 0.6518, Val. Loss: 0.6707, Val. Acc.: 0.5964  
 Training: 100%| | 140/140 [06:45<00:00, 2.90s/it, loss=0.508]  
 Epoch: 006, Train Loss: 0.6407, Val. Loss: 0.6714, Val. Acc.: 0.5803  
 Training: 100%| | 140/140 [06:51<00:00, 2.94s/it, loss=0.683]  
 Epoch: 007, Train Loss: 0.6360, Val. Loss: 0.6722, Val. Acc.: 0.5892  
 Training: 100%| | 140/140 [06:47<00:00, 2.91s/it, loss=0.805]  
 Saved Best Model: Epoch 8, Val. Acc.: 0.6233  
 Epoch: 008, Train Loss: 0.6310, Val. Loss: 0.6810, Val. Acc.: 0.6233  
 Training: 100%| | 140/140 [06:49<00:00, 2.92s/it, loss=0.801]  
 Epoch: 009, Train Loss: 0.6164, Val. Loss: 0.6860, Val. Acc.: 0.6036  
 Training: 100%| | 140/140 [06:56<00:00, 2.97s/it, loss=0.667]  
 Epoch: 010, Train Loss: 0.6062, Val. Loss: 0.6926, Val. Acc.: 0.6170  
 Training: 100%| | 140/140 [06:51<00:00, 2.94s/it, loss=0.585]  
 Epoch: 011, Train Loss: 0.6008, Val. Loss: 0.7020, Val. Acc.: 0.5982  
 Training: 100%| | 140/140 [06:48<00:00, 2.92s/it, loss=0.922]  
 Epoch: 012, Train Loss: 0.5924, Val. Loss: 0.7067, Val. Acc.: 0.5892  
 Training: 100%| | 140/140 [06:52<00:00, 2.95s/it, loss=0.719]  
 Epoch: 013, Train Loss: 0.5746, Val. Loss: 0.7608, Val. Acc.: 0.6117  
 Training: 100%| | 140/140 [06:51<00:00, 2.94s/it, loss=0.463]  
 Epoch: 014, Train Loss: 0.5567, Val. Loss: 0.7508, Val. Acc.: 0.5964  
 Training: 100%| | 140/140 [06:51<00:00, 2.94s/it, loss=0.525]  
 Epoch: 015, Train Loss: 0.5353, Val. Loss: 0.7901, Val. Acc.: 0.5830  
 Training: 100%| | 140/140 [06:47<00:00, 2.91s/it, loss=0.425]  
 Early Stopping

[ ]:

[ ]:

[ ]: